

# lab\_nlls\_soln

February 24, 2019

## 1 Lab: Nonlinear Least Squares for Modeling Materials

In this lab, we will explore gradient descent on nonlinear least squares. Suppose we wish to fit a model of the form,

$$\hat{y} = f(x, w)$$

where  $w$  is a vector of parameters and  $x$  is the vector of predictors. In nonlinear least squares, we find  $w$  by minimizing a least-squares function

$$J(w) = \sum_i (y_i - f(x_i, w))^2$$

where the summation is over training samples  $(x_i, y_i)$ . In general, this optimization has no closed-form expression. So gradient descent is widely used.

In this lab, we will implement gradient descent on nonlinear least squares in physical modeling of materials. Specifically, we will estimate parameters for expansion of copper as a function of temperature. In doing this lab, you will learn to:

- \* Set up a nonlinear least squares as an unconstrained optimization function
- \* Compute initial parameter estimates for a simple rational model
- \* Compute the gradients of the least squares objective
- \* Implement gradient descent for minimizing the objective
- \* Implement momentum gradient descent
- \* Visualize the convergence of the algorithm

We first import some key packages.

```
In [17]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.linear_model import Ridge, LinearRegression
```

### 1.1 Load the Data

The NIST agency has an excellent [nonlinear regression website](#) that has several datasets for trying nonlinear regression problem. In this lab, we will use the data from a NIST study involving the thermal expansion of copper. The response variable is the coefficient of thermal expansion, and the predictor variable is temperature in degrees kelvin.

Hahn, T., NIST (1979), Copper Thermal Expansion Study. (unpublished)

You can download the data as follows.


```
In [18]: url = 'https://itl.nist.gov/div898/strd/nls/data/LINKS/DATA/Hahn1.dat'
df = pd.read_csv(url, skiprows=60, sep=' ', skipinitialspace=True, names=['x0', 'y0', 'dummy'])
df.head()
```

```
Out[18]:
```

	x0	y0	dummy
0	0.591	24.41	NaN
1	1.547	34.82	NaN
2	2.902	44.09	NaN
3	2.894	45.07	NaN
4	4.703	54.98	NaN

Extract the x0 and y0 into arrays. Rescale, x0 and y0 to values between 0 and 1 by dividing x0 and y0 by the maximum value. Store the scaled values in vectors x and y. The rescaling will help with the conditioning of the fitting. Plot, y vs. x.

```
In [19]: # TODO
# x0 = ...
# y0 = ...
# x0 = x0/np.max(x0)
# y0 = y0/np.max(y0)
# plt.plot(...)
x0 = np.array(df['x0'])
y0 = np.array(df['y0'])
x = x0/np.max(x0)
y = y0/np.max(y0)
plt.plot(x,y, 'o')
plt.grid()
```




To make the problem a little more challenging, we will add some noise. Add random Gaussian noise with mean 0 and std. dev = 0.05 to  $y$ . Store the noisy results in  $yn$ . You can use the `np.random.normal()` function to add Gaussian noise. Plot  $yn$  vs.  $x$ .

```
In [20]: # TODO
        #  $yn = y + \dots$ 
        n = len(y0)
        yn = y + np.random.normal(0,0.02,n)

        plt.plot(x,y,'o')
        plt.plot(x,yn,'o')

Out[20]: [<matplotlib.lines.Line2D at 0x28edbc31a58>]
```



Split the data  $(x, yn)$  into training and test. Let  $x_{tr}, y_{tr}$  be training data and  $x_{ts}, y_{ts}$  be the test data. You can use the `train_test_split` function. Set `test_size=0.33` so that 1/3 of the samples are held out for test.

```
In [21]: from sklearn.model_selection import train_test_split

        # TODO
        #  $x_{tr}, x_{ts}, y_{tr}, y_{ts} = \dots$ 
        xtr, xts, ytr, yts = train_test_split(x,yn,test_size=0.33)
```

## 1.2 Initial Fit for a Rational Model

The [NIST website](#) suggests using a *rational* model of the form,

$$\hat{y} = (a[0] + a[1]*x + \dots + a[d]*x^d)/(1 + b[0]*x + \dots + b[d-1]*x^d)$$

with  $d=3$ . The model parameters are  $w = [a[0], \dots, a[d], b[0], \dots, b[d-1]]$  so there are  $2d+1$  parameters total. Complete the function below that takes vectors  $w$  and  $x$  and predicts a set of values  $\hat{y}$  using the above model.

```
In [22]: def predict(w,x):
```

```
    # Get the length
    d = (len(w)-1)//2

    # TODO. Extract a and b from w
    # a = ...
    # b = ...
    a = w[:d+1]
    b = w[d+1:]

    # TODO. Compute yhat. You may use the np.polyval function
    # But, remember you must flip the order the a and b
    # yhat = ...
    arev = a[::-1]
    brev = b[::-1]
    znum = np.polyval(arev,x)
    zden = 1+x*np.polyval(brev,x)
    yhat = znum/zden
    return yhat
```

When we fit with a nonlinear model, most methods only get convergence to a local minima. So, you need a good initial condition. For a rational model, one way to get is to realize that if:

$$y \approx (a[0] + a[1]*x + \dots + a[d]*x^d)/(1 + b[0]*x + \dots + b[d-1]*x^d)$$

Then:

$$y \approx a[0] + a[1]*x + \dots + a[d]*x^d - b[0]*x*y + \dots - b[d-1]*x^d*y.$$

So, we can solve for the parameters  $w = [a, b]$  from linear regression of the predictors,

$$Z[i,:] = [x[i], \dots, x[i]**d, y[i]*x[i], \dots, y[i]*x[i]**d]$$

```
In [23]: d = 3
```

```
    # TODO. Create the transformed feature matrix
    # Z = ...
    powd = np.arange(1,d+1)
```

```

Znum = xtr[:,None]**powd[None,:]
Zden = -ytr[:,None]*Znum
Z = np.hstack((Znum, Zden))

# TODO. Fit with parameters with linear regression
# regr = LinearRegression()
# regr.fit(...)
regr = LinearRegression()
regr.fit(Z,ytr)

# TODO
# Extract the parameters from regr.coef_ and regr.intercept_ and store the parameter
# winit = ...
w0 = regr.coef_
a0 = regr.intercept_
winit = np.hstack((a0,w0))

```

Now plot the predicted values of the  $\hat{y}$  vs.  $x$  using your estimated parameter `winit` for 1000 values  $x$  in  $[0, 1]$ . On the same plot, plot  $y$  vs.  $x$ . You will see that you get a horrible fit.

```


In [24]: # TODO
# xp = ...
# yhat = ...
# plot(...)
xp = np.linspace(0,1,1000)
yhat = predict(winit,xp)
plt.plot(xp,yhat)
plt.plot(xts, yts, 'o')

```

```

Out[24]: [<matplotlib.lines.Line2D at 0x28edbca7828>]

```




The reason the previous fit is poor is that the denominator in  $\hat{y}$  goes close to zero. To avoid this problem, we can use Ridge regression, to try to keep the parameters close to zero. Re-run the fit above with Ridge with  $\alpha = 1e-3$ . You should see you get a reasonable, but not perfect fit.

```
In [25]: # TODO. Fit with parameters with linear regression
# regr = Ridge(alpha=1e-3)
# regr.fit(...)
regr = Ridge(alpha=1e-3)
regr.fit(Z,ytr)

# TODO
# Extract the parameters from regr.coef_ and regr.intercept_
# winit = ...
w0 = regr.coef_
a0 = regr.intercept_
winit = np.hstack((a0,w0))

# TODO
# Plot the results as above.
xp = np.linspace(0,1,1000)
yhat = predict(winit,xp)
plt.plot(xp,yhat)
plt.plot(xts, yts, 'o')
```

```
Out[25]: [<matplotlib.lines.Line2D at 0x28edbd074a8>]
```



### 1.3 Creating a Loss Function

We can now use gradient descent to improve our initial estimate. Complete the following function to compute

$f(w) = 0.5 * \sum_i (y[i] - \hat{y}[i])^2$

and fgrad, the gradient of  $f(w)$ .

```
In [26]: def feval(w,x,y):

    # TODO. Parse w
    # a = ...
    # b = ...
    d = (len(w)-1)//2
    a = w[:d+1]
    b = w[d+1:]

    # TODO. Znum[i,j] = x[i]**j
    pow1 = np.arange(0,d+1)
    Znum = x[:,None]**pow1[None,:]

    # TODO. Zden[i,j] = x[i]**(j+1)
    pow2 = np.arange(1,d+1)
    Zden = x[:,None]**pow2[None,:]

    # TODO. Compute yhat
    # Compute the numerator and denominator
    rnum = Znum.dot(a)
    rden = Zden.dot(b)
    yhat = rnum/(1+rden)

    # TODO. Compute loss
    # f = ...
    e = yhat-y
    f = 0.5*np.sum(e**2)

    # TODO. Compute gradients
    # fgrad = ...
    eden = e/(1+rden)
    dJ_da = eden.dot(Znum)
    enum = -e*yhat/(1+rden)
    dJ_db = enum.dot(Zden)
    fgrad = np.hstack((dJ_da, dJ_db))

    return f, fgrad
```

Test the gradient function: \* Take  $w_0=w_{init}$  and compute  $f_0, f_{grad0} = \text{feval}(w_0, x_{tr}, y_{tr})$   
\* Take  $w_1$  very close to  $w_0$  and compute  $f_1, f_{grad1} = \text{feval}(w_1, x_{tr}, y_{tr})$  \* Verify that  $f_1-f_0$  is close to the predicted value based on the gradient.

```
In [27]: # TODO
w0 = winit
p = len(winit)
w1 = w0 + np.random.normal(0,1,p)*1e-6
f0, fgrad0 = feval(w0,xtr,ytr)
f1, fgrad1 = feval(w1,xtr,ytr)

print([f1-f0, fgrad0.dot(w1-w0)])
```

```
[7.116281790800483e-06, 7.1106249555072314e-06]
```

## 1.4 Implement gradient descent


We will now try to minimize the loss function with gradient descent. Using the function `feval` defined above, implement gradient descent. Run gradient descent with a step size of  $\alpha=1e-6$  starting at  $w=w_{init}$ . Run it for  $nit=10000$  iterations. Compute  $fgd[it]$ = the objective function on iteration  $it$ . Plot  $fgd[it]$  vs.  $it$ .

You should see that the training loss decreases, but it still hasn't converged after 10000 iterations.

```
In [28]: nit = 10000
step = 1e-6
wt = winit
fgd = np.zeros(nit)
for it in range(nit):
    ft, fgradt = feval(wt,xtr,ytr)
    fgd[it] = ft
    wt = wt - step*fgradt

plt.plot(fgd)
plt.xlabel('Iteration')
plt.ylabel('Training loss')
plt.grid()
```





Now, try to get a faster convergence with adaptive step-size using the Armijo rule. Implement the gradient descent with adaptive step size. Let  $f_{\text{adapt}}[it]$  be the loss function on iteration  $it$ . Plot  $f_{\text{adapt}}[it]$  and  $f_{\text{gd}}[it]$  vs.  $it$  on the same graph. You should see a slight improvement, but not much.


```
In [29]: nit = 10000
         step = 1e-6 # Initial step
         w0 = winit
         fadapt = np.zeros(nit)
         f0, fgrad0 = feval(w0,xtr,ytr)
         for it in range(nit):

             # Compute test point
             w1 = w0 - step*fgrad0
             f1, fgrad1 = feval(w1,xtr,ytr)

             # Test Armijo rule
             alpha = 0.5
             if (f1-f0 < alpha*fgrad0.dot(w1-w0)) and (f1 < f0):
                 step = step*2
                 f0 = f1
                 fgrad0 = fgrad1
                 w0 = w1
             else:
                 step = 0.5*step
```

```
fadapt[it] = f0
```


```
plt.plot(fgd)
plt.plot(fadapt)
plt.xlabel('Iteration')
plt.ylabel('Training loss')
plt.legend(['Fixed step', 'Adaptive step'])
plt.grid()
```



Using the final estimate for  $w$  from the adaptive step-size plot the predicted values of the  $yhat$  vs.  $x$  for 1000 values  $x$  in  $[0,1]$ . On the same plot, plot  $yhat$  vs.  $x$  for the initial parameter  $w=winit$ . Also, plot  $yts$  vs.  $xts$ . You should see that gradient descent was able to improve the estimate slightly, although the initial estimate was not too bad.

```
In [30]: xp = np.linspace(0,1,500)
         yinit = predict(winit,xp)
         yhat = predict(w0,xp)
         plt.plot(xts,yts, 'o')
         plt.plot(xp,yinit, '-')
         plt.plot(xp,yhat, '-')
```

```
Out[30]: [<matplotlib.lines.Line2D at 0x28edce39978>]
```



## 1.5 Momentum Gradient Descent

This section is bonus.

One way to improve gradient descent is to use *momentum*. In momentum gradient descent, the update rule is:

```
f, fgrad = feval(w,...)
z = beta*d + fgrad
w = w - step*z
```


This is similar to gradient descent, except that there is a second order term on the gradient. Implement this algorithm with  $\beta = 0.99$  and  $\text{step} = 1e-5$ . Compare the convergence of the loss function with gradient descent.

```
In [31]: nit = 1000
         step = 1e-5
         beta = 0.99
         wt = winit
         p = len(winit)
         z = np.zeros(p)
         fmom = np.zeros(nit)
         for it in range(nit):
             ft, fgradt = feval(wt,xtr,ytr)
             z = beta*z + fgradt
             wt = wt - step*z
             fmom[it] = ft
```

```

plt.plot(fgd)
plt.plot(fadapt)
plt.plot(fmom)
plt.xlabel('Iteration')
plt.ylabel('Training loss')
plt.legend(['Fixed step', 'Adaptive step', 'Momentum'])
plt.grid()

```



```


In [32]: xp = np.linspace(0,1,500)
         yinit = predict(winit,xp)
         yhat = predict(wt,xp)
         plt.plot(xts,yts,'o')
         plt.plot(xp,yinit,'-')
         plt.plot(xp,yhat,'-')

```

```

Out[32]: [<matplotlib.lines.Line2D at 0x28edcf15c88>]

```



In [ ]:

In [ ]: